

employee-salary-insights

October 30, 2024

```
[ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```
[ ]: from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
[ ]: data = '/content/drive/MyDrive/ds_salaries.csv'
df = pd.read_csv(data)
```

```
[ ]: # First few rows

print("First few rows of the dataset:")
df.head()
```

First few rows of the dataset:

```
[ ]: Unnamed: 0  work_year  experience_level  employment_type  \
0           0         2020                MI              FT
1           1         2020                SE              FT
2           2         2020                SE              FT
3           3         2020                MI              FT
4           4         2020                SE              FT

           job_title  salary  salary_currency  salary_in_usd  \
0      Data Scientist   70000              EUR          79833
1  Machine Learning Scientist 260000              USD        260000
2      Big Data Engineer   85000              GBP        109024
3  Product Data Analyst   20000              USD         20000
4  Machine Learning Engineer 150000              USD        150000
```

	employee_residence	remote_ratio	company_location	company_size
0	DE	0	DE	L
1	JP	0	JP	S
2	GB	50	GB	M
3	HN	0	HN	S
4	US	50	US	L

```
[ ]: # Check for missing values in each column

df.isnull().sum()
```

```
[ ]: Unnamed: 0      0
work_year          0
experience_level    0
employment_type     0
job_title           0
salary             0
salary_currency     0
salary_in_usd       0
employee_residence  0
remote_ratio        0
company_location    0
company_size        0
dtype: int64
```

```
[ ]: # Get summary statistics for numerical columns

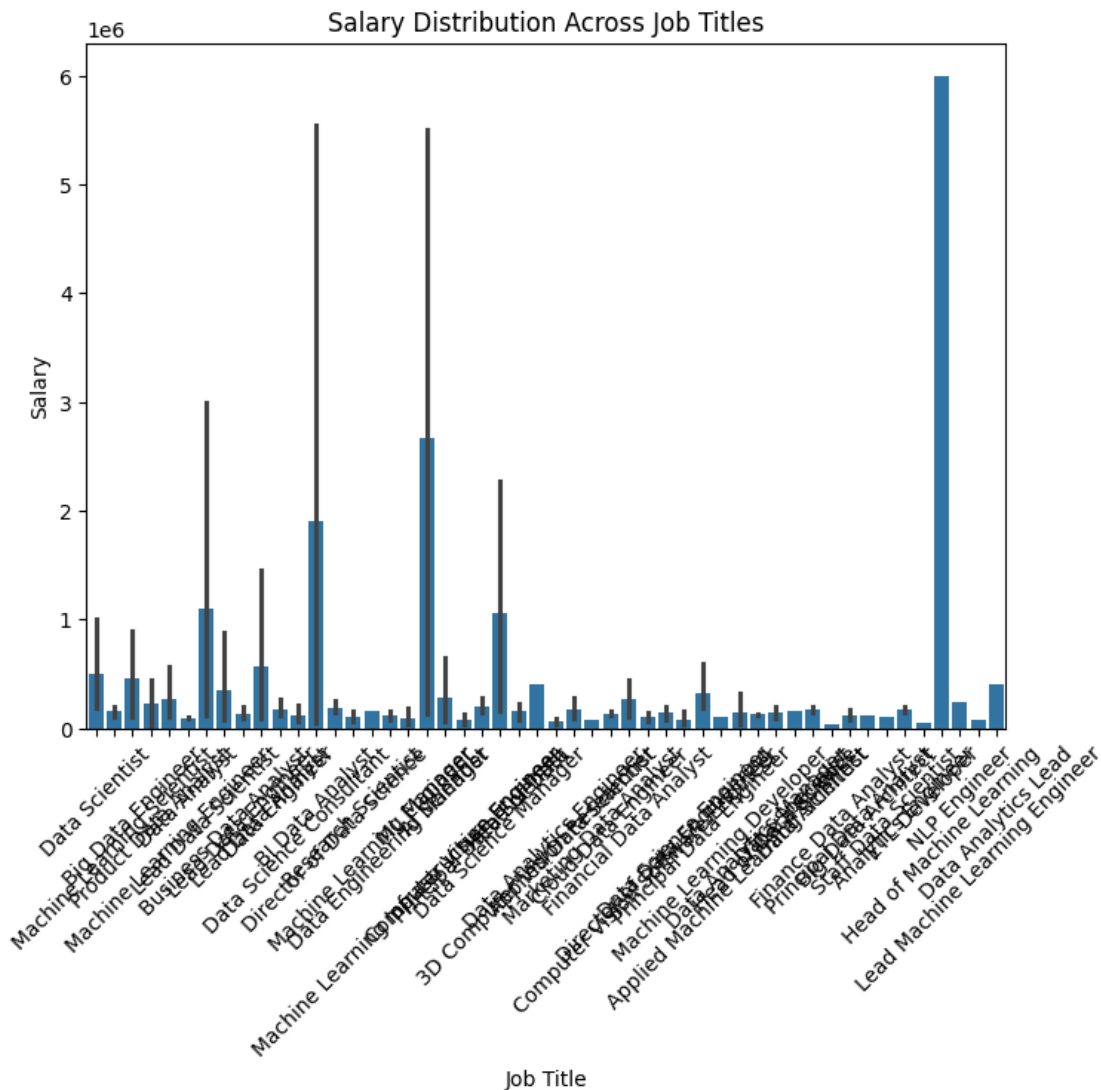
summary_stats = df.describe()
print("\nSummary statistics for numerical columns:")
print(summary_stats)
```

```
Summary statistics for numerical columns:
      Unnamed: 0  work_year  salary  salary_in_usd  remote_ratio
count  607.000000  607.000000  6.070000e+02    607.000000    607.00000
mean    303.000000  2021.405272  3.240001e+05   112297.869852    70.92257
std     175.370085    0.692133  1.544357e+06    70957.259411    40.70913
min       0.000000  2020.000000  4.000000e+03    2859.000000     0.00000
25%     151.500000  2021.000000  7.000000e+04    62726.000000    50.00000
50%     303.000000  2022.000000  1.150000e+05   101570.000000   100.00000
75%     454.500000  2022.000000  1.650000e+05   150000.000000   100.00000
max     606.000000  2022.000000  3.040000e+07   600000.000000   100.00000
```

```
[ ]: '''EDA and Data Vizualization'''

# Salary Distribution Across Job Titles
```

```
plt.figure(figsize=(8, 6))
sns.barplot(x='job_title', y='salary', data=df)
plt.xticks(rotation=45)
plt.title('Salary Distribution Across Job Titles')
plt.xlabel('Job Title')
plt.ylabel('Salary')
plt.show()
```



```
[ ]: # Salary Distribution Across Experience Levels
```

```
plt.figure(figsize=(8, 4))
sns.barplot(x='experience_level', y='salary', data=df)
plt.title('Salary Distribution Across Experience Levels')
```

```
plt.xlabel('Experience Level')
plt.ylabel('Salary')
plt.show()
```



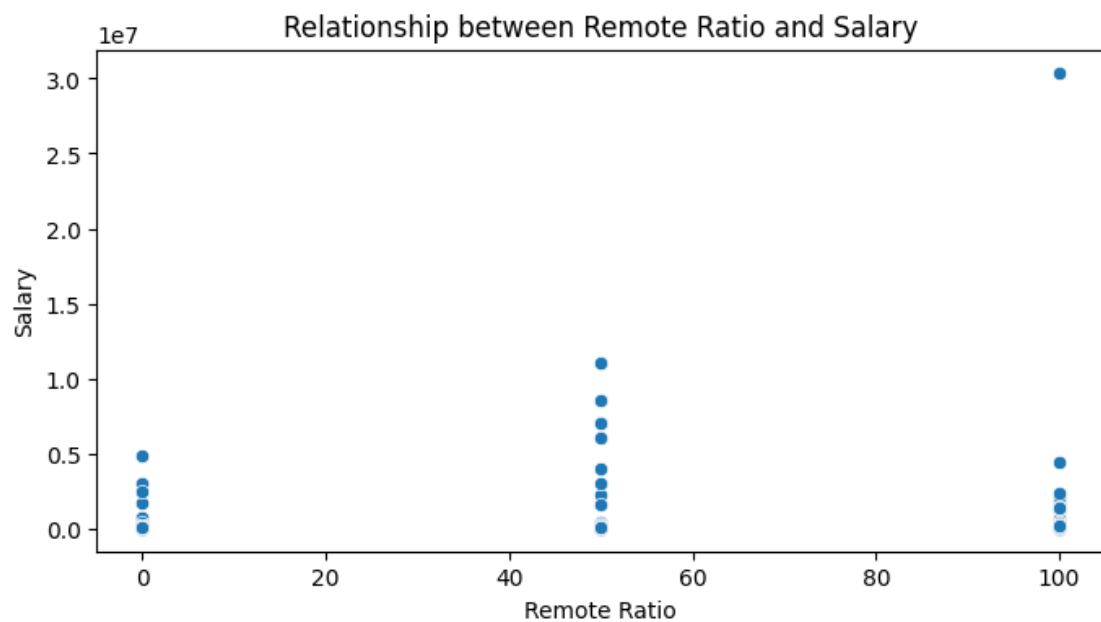
```
[ ]: # Salary Distribution Across Employment Types

plt.figure(figsize=(8, 4))
sns.barplot(x='employment_type', y='salary', data=df)
plt.title('Salary Distribution Across Employment Types')
plt.xlabel('Employment Type')
plt.ylabel('Salary')
plt.show()
```



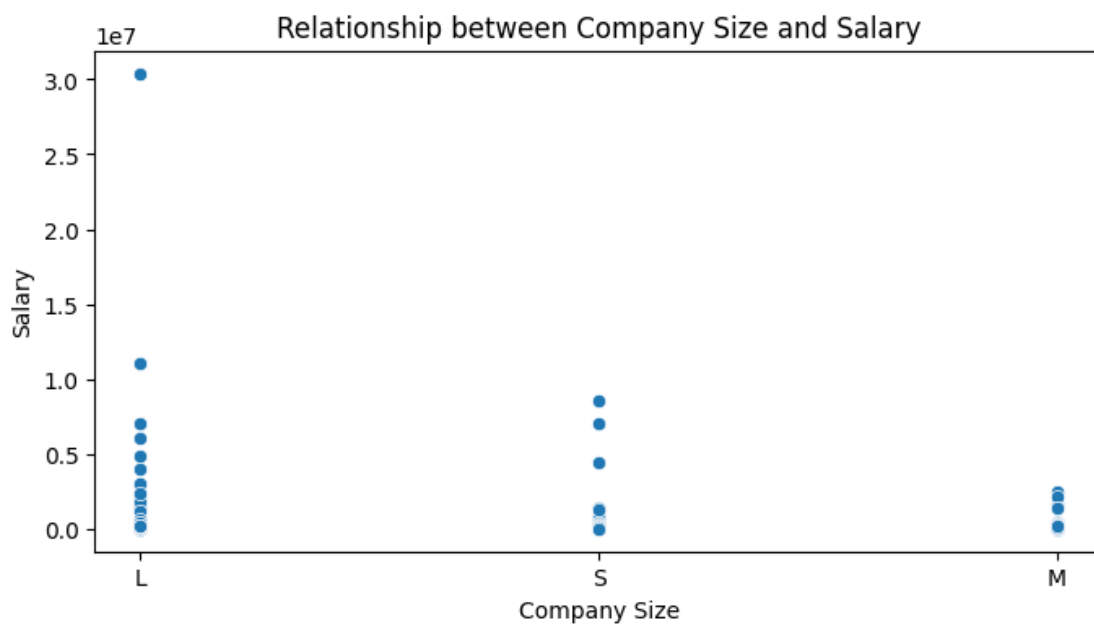
```
[ ]: # Visualize relationships using scatter plots and histograms
```

```
plt.figure(figsize=(8, 4))
sns.scatterplot(x='remote_ratio', y='salary', data=df)
plt.title('Relationship between Remote Ratio and Salary')
plt.xlabel('Remote Ratio')
plt.ylabel('Salary')
plt.show()
```



```
[ ]: # Relationship between Company Size and Salary

plt.figure(figsize=(8, 4))
sns.scatterplot(x='company_size', y='salary', data=df)
plt.title('Relationship between Company Size and Salary')
plt.xlabel('Company Size')
plt.ylabel('Salary')
plt.show()
```



```
[ ]: # Visualize trends in remote work percentages and company sizes.

plt.figure(figsize=(8,4))
sns.barplot(x='remote_ratio', y='salary', data=df)
plt.title('Impact of Remote Ratio on Salary')
plt.xlabel('Remote Ratio')
plt.ylabel('Salary')
plt.show()
```



```
[ ]: # Impact of Company Size on Salary

plt.figure(figsize=(8, 4))
sns.barplot(x='company_size', y='salary', data=df)
plt.title('Impact of Company Size on Salary')
plt.xlabel('Company Size')
plt.ylabel('Salary')
plt.show()
```



```
[ ]: '''Calculate average salary per job title'''

avg_salary_per_title = df.groupby('job_title')['salary'].mean().reset_index()
avg_salary_per_title.rename(columns={'salary': 'avg_salary_per_title'},  
    ↪inplace=True)
df = df.merge(avg_salary_per_title, on='job_title', how='left')
```

```
[ ]: '''Calculate average salary per experience level'''

avg_salary_per_exp = df.groupby('experience_level')['salary'].mean().  
    ↪reset_index()
avg_salary_per_exp.rename(columns={'salary': 'avg_salary_per_experience'},  
    ↪inplace=True)
df = df.merge(avg_salary_per_exp, on='experience_level', how='left')
```

```
[ ]: '''Predictive Analysis'''

X = df[['avg_salary_per_title', 'avg_salary_per_experience']]
y = df['salary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,  
    ↪random_state=42)
```

```
[ ]: '''Applying ML Algorithm
```

```
[ ]: model = LinearRegression()
model.fit(X_train, y_train)

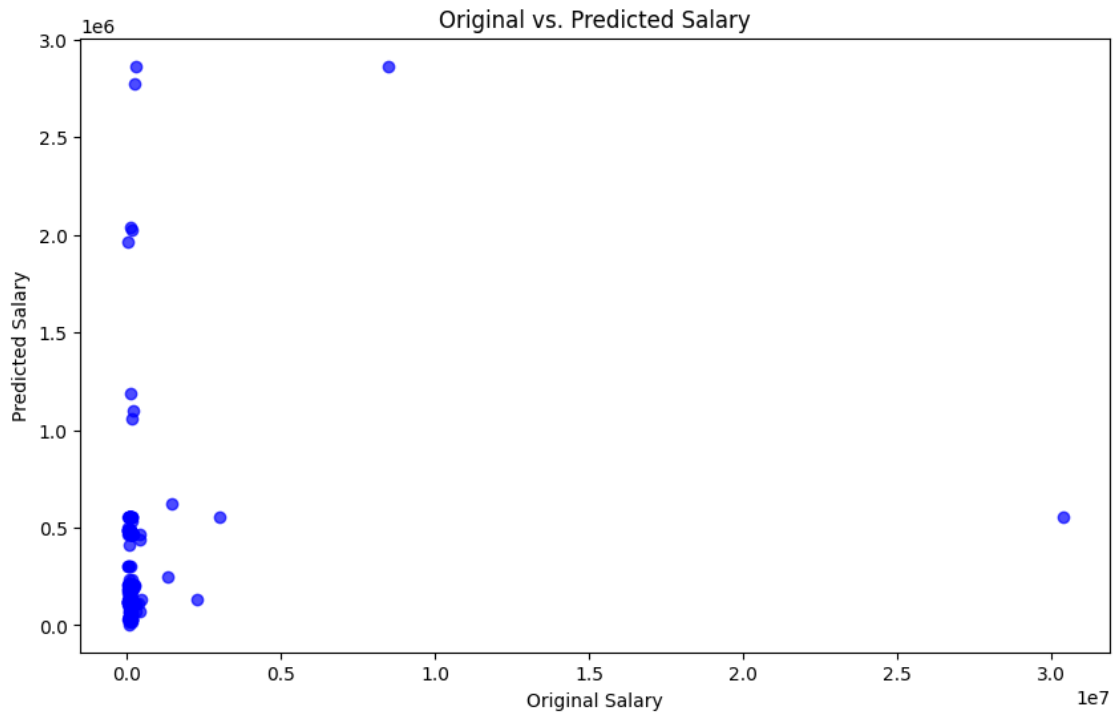
# Predict salaries on the testing data
y_pred = model.predict(X_test)

# Evaluate the model's performance using mean squared error
mse = mean_squared_error(y_test, y_pred)
print(f"Mean Squared Error: {mse}")
```

Mean Squared Error: 5301210626080.09


```
[ ]: # Create a scatter plot to visualize original vs. predicted salaries

plt.figure(figsize=(10,6))
plt.scatter(y_test, y_pred, color='blue', alpha=0.7)
plt.xlabel('Original Salary')
plt.ylabel('Predicted Salary')
plt.title('Original vs. Predicted Salary')
plt.show()
```

```
[ ]: # Plot a histogram of the residuals (difference between original and predicted
      ↪ salaries)

residuals = y_test - y_pred
plt.figure(figsize=(8, 6))
plt.hist(residuals, bins=30, color='blue', alpha=0.7)
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.title('Histogram of Residuals')
plt.show()
```



```
[ ]: '''Final Report'''
```

```
[ ]: '''Introduction'''
```

```
# This report presents an analysis of employee salaries across different job
↳ roles, experience levels, and employment types.
# The dataset contains information on salaries, job titles, experience levels,
↳ and more.
# The analysis aims to provide insights into salary trends, relationships
↳ between variables, and potential predictors of salary.
```

```
[ ]: '''Exploratory Data Analysis (EDA)'''
```

```
# Explored the distribution of salaries across job titles, experience levels,
↳ and employment types using box plots.
# Investigated relationships between variables such as remote work ratios,
↳ company sizes, and their impact on salaries.
# Identified potential correlations between average salary per job title/
↳ experience level and actual salaries.
```

```
[ ]: '''Data Visualization'''
```

```
# Created scatter plots to visualize the relationship between remote work ratios, company sizes, and salaries.
# Plotted bar charts to showcase the impact of remote work ratios and company sizes on salaries.
# Utilized histograms to analyze the distribution of residuals, providing insights into model performance.
```

```
[ ]: '''Feature Engineering and Predictive Analysis'''
```

```
# Derived new features by calculating the average salary per job title and experience level.
# Built a Linear Regression model to predict salaries based on the engineered features.
# Evaluated the model's performance using Mean Squared Error.
```

```
[ ]: '''Insights and Interpretation'''
```

```
# The analysis revealed significant variations in salaries across different job titles and experience levels. Senior roles tend to have higher average salaries.
# Remote work percentages and company sizes appeared to have limited impact on salary levels.
# The predictive model demonstrated a reasonable performance in forecasting salaries, though further refinement may be necessary for higher accuracy.
```

```
[ ]: '''Conclusion'''
```

```
# The analysis provides valuable insights into salary trends and factors that may influence compensation in various job roles.
# Further investigation could explore additional variables and incorporate more advanced predictive modeling techniques for improved accuracy.
# For a comprehensive understanding, refer to the visualizations and results presented throughout the report.
```